Stanford CS 229, Public Course, Problem Set 1

Dylan Price

December 2, 2016

1

a)

Find the Hessian of the cost function $J(\theta) = \frac{1}{2}\sum_{i=1}^{m}(\theta^T x^{(i)} - y^{(i)})^2$

We know that $H_{jk} = \frac{\partial^2 J(\theta)}{\partial\theta_j\theta_k}$

First find $\frac{\partial J(\theta)}{\partial\theta_k}$,

$$\begin{aligned}\frac{\partial J(\theta)}{\partial\theta_k} &= \frac{1}{2}\sum_{i=1}^{m}\frac{\partial}{\partial\theta_k}(\theta^T x^{(i)} - y^{(i)})^2 \\ &= \frac{1}{2}\sum_{i=1}^{m}2(\theta^T x^{(i)} - y^{(i)})(x_k^{(i)}) \\ &= \sum_{i=1}^{m}(\theta^T x^{(i)} - y^{(i)})(x_k^{(i)})\end{aligned}$$

Now find $\frac{\partial^2 J(\theta)}{\partial\theta_j\theta_k}$,

$$\begin{aligned}\frac{\partial^2 J(\theta)}{\partial\theta_j\theta_k} &= \frac{\partial}{\partial\theta_j}(\frac{\partial J(\theta)}{\partial\theta_k}) \\ &= \sum_{i=1}^{m}\frac{\partial}{\partial\theta_j}((\theta^T x^{(i)} - y^{(i)})(x_k^{(i)})) \\ &= \sum_{i=1}^{m}x_j^{(i)}x_k^{(i)} \text{ for } 1 \le j \le n \text{ and } 1 \le k \le n\end{aligned}$$

Therefore

$$
\begin{aligned}
H &= \begin{bmatrix}
\sum_{i=1}^{m} x_1^{(i)} x_1^{(i)} & \sum_{i=1}^{m} x_2^{(i)} x_1^{(i)} & \cdots & \sum_{i=1}^{m} x_n^{(i)} x_1^{(i)} \\
\sum_{i=1}^{m} x_1^{(i)} x_2^{(i)} & \sum_{i=1}^{m} x_2^{(i)} x_2^{(i)} & \cdots & \sum_{i=1}^{m} x_n^{(i)} x_2^{(i)} \\
\vdots & \vdots & \ddots & \vdots \\
\sum_{i=1}^{m} x_1^{(i)} x_n^{(i)} & \sum_{i=1}^{m} x_2^{(i)} x_n^{(i)} & \cdots & \sum_{i=1}^{m} x_n^{(i)} x_n^{(i)}
\end{bmatrix} \\
&= \begin{bmatrix} x_1^{(1)} & \cdots & x_1^{(n)} \\ \vdots & \ddots & \vdots \\ x_m^{(1)} & \cdots & x_m^{(n)} \end{bmatrix} \begin{bmatrix} x_1^{(1)} & \cdots & x_m^{(1)} \\ \vdots & \ddots & \vdots \\ x_1^{(n)} & \cdots & x_m^{(n)} \end{bmatrix} \\
&= X^T X
\end{aligned}
$$

b)

Show that the first iteration of Newton's method gives us $\theta^* = (X^T X)^{-1} X^T \vec{y}$, the solution to our least squares problem.

One iteration of Newton's Method:
$\theta := \theta - H^{-1} \nabla_\theta J(\theta)$
Therefore,

$$
\begin{aligned}
\theta^* &= \theta - (X^T X)^{-1} \begin{bmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_m} \end{bmatrix} \\
&= \theta - (X^T X)^{-1} \begin{bmatrix} \sum_{i=1}^{m} (x_1^{(i)} \theta^T x^{(i)} - x_1^{(i)} y^{(i)}) \\ \vdots \\ \sum_{i=1}^{m} (x_m^{(i)} \theta^T x^{(i)} - x_m^{(i)} y^{(i)}) \end{bmatrix}
\end{aligned}
$$

let $\theta = \vec{0}$ (initialize θ)

$$
\begin{aligned}
&= -(X^T X)^{-1} \begin{bmatrix} \sum_{i=1}^{m} (-x_1^{(i)} y^{(i)}) \\ \vdots \\ \sum_{i=1}^{m} (-x_m^{(i)} y^{(i)}) \end{bmatrix} \\
&= (X^T X)^{-1} \begin{bmatrix} \sum_{i=1}^{m} (x_1^{(i)} y^{(i)}) \\ \vdots \\ \sum_{i=1}^{m} (x_m^{(i)} y^{(i)}) \end{bmatrix} \\
&= (X^T X)^{-1} X^T \vec{y}
\end{aligned}
$$

2

a

See q2/ folder

b

At low values of τ, the classification boundaries are clustered around the positive training examples. As you increase τ, these boundaries begin to merge into bigger areas, i.e. the classification boundaries look less 'local' to the positive training examples. At high values of τ, the classification boundary is essentially a straight line dividing positive and negative classes.

The decision boundary of unweighted logistic regression would look like the plots with the highest values of τ. This is because as τ approaches infinity, the weight function $w^{(i)}$ goes to 1 for every training example, making the regression unweighted (that is, every training example gets the same weight).

3

a)

$$\begin{aligned} J(\theta) &= \frac{1}{2}\sum_{i=1}^{m}\sum_{j=1}^{p}((\theta^T x^{(i)})_j - y_j^{(i)})^2 \\ &= \frac{1}{2}\sum_{i=1}^{m}(\theta^T x^{(i)} - y^{(i)})^T(\theta^T x^{(i)} - y^{(i)}) \\ &\text{Let } \vec{1} \text{ be a vector of all ones.} \\ &= \frac{1}{2}\vec{1}^T(X\theta - Y)^T(X\theta - Y)\vec{1} \end{aligned}$$

b)

$$\begin{aligned} J(\theta) &= \frac{1}{2}\vec{1}^T(X\theta - Y)^T(X\theta - Y)\vec{1} \\ &= \frac{1}{2}\vec{1}^T((X\theta)^T - Y^T)(X\theta - Y)\vec{1} \\ &= \frac{1}{2}\vec{1}^T(\theta^T X^T - Y^T)(X\theta - Y)\vec{1} \\ &= \frac{1}{2}\vec{1}^T(\theta^T X^T X\theta - \theta^T X^T Y - Y^T X\theta + Y^T Y)\vec{1} \\ &= \frac{1}{2}\vec{1}^T\theta^T X^T X\theta\vec{1} - \frac{1}{2}\vec{1}^T\theta^T X^T Y\vec{1} - \frac{1}{2}\vec{1}^T Y^T X\theta\vec{1} + \frac{1}{2}\vec{1}^T Y^T Y\vec{1} \end{aligned}$$

$$\begin{aligned} \Delta_\theta J(\theta) &= \frac{1}{2} * 2X^T X\theta - \frac{1}{2}X^T Y - \frac{1}{2}(Y^T X)^T \\ &= X^T X\theta - X^T Y \end{aligned}$$

$$
\begin{aligned}
0 &= X^T X\theta - X^T Y \\
X^T X\theta &= X^T Y \\
\theta &= (X^T X)^{-1} X^T Y
\end{aligned}
$$

c)

Consider the multivariate model $y^{(i)} = \theta^T x^{(i)}$

$$
y^{(i)} = \theta^T x^{(i)} = \begin{bmatrix} \theta_{1,1} & \theta_{1,2} & \cdots & \theta_{1,n} \\ \theta_{2,1} & \theta_{2,2} & & \\ \vdots & & \ddots & \\ \theta_{p,1} & & & \theta_{p,n} \end{bmatrix} \begin{bmatrix} x_1^{(i)} \\ x_2^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} = \begin{bmatrix} \theta_{1,1}x_1^{(i)} & \theta_{1,2}x_2^{(i)} & \cdots & \theta_{1,n}x_n^{(i)} \\ \theta_{2,1}x_1^{(i)} & \theta_{2,2}x_2^{(i)} & & \\ \vdots & & \ddots & \\ \theta_{p,1}x_1^{(i)} & & & \theta_{p,n}x_n^{(i)} \end{bmatrix} = \begin{bmatrix} \theta_1^T x^{(i)} \\ \theta_2^T x^{(i)} \\ \vdots \\ \theta_p^T x^{(i)} \end{bmatrix} = \begin{bmatrix} y_1^{(i)} \\ y_2^{(i)} \\ \vdots \\ y_p^{(i)} \end{bmatrix}
$$

Therefore each row of this multivariate model corresponds to an equation of the form $y_j^{(i)} = \theta_j^T x^{(i)}$.

4

a)

$$
\begin{aligned}
\ell(\varphi) &= \log \prod_{i=1}^{m} P(x^{(i)}, y^{(i)}; \varphi) \\
&= \log \prod_{i=1}^{m} P(y^{(i)}) P(x^{(i)} | y^{(i)}) \\
&= \sum_{i=1}^{m} \log P(y^{(i)}) P(x^{(i)} | y^{(i)}) \\
&= \sum_{i=1}^{m} [\log P(y^{(i)}) + \log \prod_{j=1}^{n} (\phi_{j|y^{(i)}})^{x_j^{(i)}} (1 - \phi_{j|y^{(i)}})^{1 - x_j^{(i)}}] \\
&= \sum_{i=1}^{m} [\log P(y^{(i)}) + \sum_{j=1}^{n} (\log(\phi_{j|y^{(i)}})^{x_j^{(i)}} + \log(1 - \phi_{j|y^{(i)}})^{1 - x_j^{(i)}})] \\
&= \sum_{i=1}^{m} [\log P(y^{(i)}) + \sum_{j=1}^{n} (x_j^{(i)} \log(\phi_{j|y^{(i)}}) + (1 - x_j^{(i)}) \log(1 - \phi_{j|y^{(i)}}))] \\
&= \sum_{i=1}^{m} [(y^{(i)} \log \phi_y + (1 - y^{(i)}) \log(1 - \phi_y) + \sum_{j=1}^{n} (x_j^{(i)} \log(\phi_{j|y^{(i)}}) + (1 - x_j^{(i)}) \log(1 - \phi_{j|y^{(i)}}))]
\end{aligned}
$$

b)

First find ϕ_y

$$
\nabla_{\phi_y} \ell(\varphi) = \sum_{i=1}^{m} [y^{(i)} \frac{1}{\phi_y} - (1 - y^{(i)}) \frac{1}{1 - \phi_y}]
$$

$$
\begin{aligned}
0 &= \sum_{i=1}^{m} [\frac{y^{(i)}}{\phi_y} - \frac{1-y^{(i)}}{1-\phi_y}] \\
&= \sum_{i=1}^{m} [y^{(i)}(1-\phi_y) - (1-y^{(i)})\phi_y] \\
&= \sum_{i=1}^{m} [y^{(i)} - y^{(i)}\phi_y - (\phi_y - y^{(i)}\phi_y)] \\
&= \sum_{i=1}^{m} [y^{(i)} - y^{(i)}\phi_y - \phi_y + y^{(i)}\phi_y] \\
&= \sum_{i=1}^{m} [y^{(i)} - \phi_y] \\
\sum_{i=1}^{m} \phi_y &= \sum_{i=1}^{m} \phi_y \\
m\phi_y &= \sum_{i=1}^{m} y^{(i)} \\
\phi_y &= \frac{\sum_{i=1}^{m} y^{(i)}}{m} = \sum_{i=1}^{m} \frac{1\{y^{(i)} = 1\}}{m}
\end{aligned}
$$

Now find $\phi_{j|y=0}$ and $\phi_{j|y=1}$

$$
\begin{aligned}
\nabla_{\phi_{j|y^{(i)}}} \ell(\varphi) = &\nabla_{\phi_{j|y^{(i)}}} \sum_{i=1}^{m} [(y^{(i)} \log \phi_y + (1-y^{(i)}) \log(1-\phi_y)] + \\
&\nabla_{\phi_{j|y^{(i)}}} \sum_{i=1}^{m} [\sum_{j=1}^{n} [(x_j^{(i)} \log(\phi_{j|y^{(i)}}) + (1-x_j^{(i)}) \log(1-\phi_{j|y^{(i)}}))]]
\end{aligned}
$$

We can drop the first term (the one containing references to ϕ_y) since its gradient with respect to $\phi_{j|y^{(i)}}$ is 0. We can also drop $\sum_{j=1}^{n}$ from the second term because we are taking the gradient with respect to a single $\phi_{j|y^{(i)}}$.

$$
\begin{aligned}
&= \nabla_{\phi_{j|y^{(i)}}} \sum_{i=1}^{m} [x^{(i)} \log \phi_{j|y^{(i)}} + (1-x^{(i)}) \log(1-\phi_{j|y^{(i)}})] \\
&= \sum_{i=1}^{m} [x^{(i)} \frac{1}{\phi_{j|y^{(i)}}} - (1-x^{(i)}) \frac{1}{1-\phi_{j|y^{(i)}}}]
\end{aligned}
$$

$$
\begin{aligned}
0 &= \sum_{i=1}^{m}[x^{(i)}\frac{1}{\phi_{j|y^{(i)}}} - (1-x^{(i)})\frac{1}{1-\phi_{j|y^{(i)}}}] \\
&= \sum_{i=1}^{m}[x^{(i)}(1-\phi_{j|y^{(i)}}) - (1-x^{(i)})\phi_{j|y^{(i)}}] \\
&= \sum_{i=1}^{m}[x^{(i)} - \phi_{j|y^{(i)}}] \\
\sum_{i=1}^{m}\phi_{j|y^{(i)}} &= \sum_{i=1}^{m}x^{(i)} \\
\sum_{i=1}^{m}\phi_{j|y=0}1\{y^{(i)}=0\} &= \sum_{i=1}^{m}x^{(i)}1\{y^{(i)}=0\} \\
\phi_{j|y=0} &= \frac{\sum_{i=1}^{m}x^{(i)}1\{y^{(i)}=0\}}{\sum_{i=1}^{m}1\{y^{(i)}=0\}} = \frac{\sum_{i=1}^{m}1\{x^{(i)}=1 \wedge y^{(i)}=0\}}{\sum_{i=1}^{m}1\{y^{(i)}=0\}}
\end{aligned}
$$

Similarly,

$$
\phi_{j|y=1} = \frac{\sum_{i=1}^{m}1\{x^{(i)}=1 \wedge y^{(i)}=1\}}{\sum_{i=1}^{m}1\{y^{(i)}=1\}}
$$

c)

We predict $y=1$ when $P(y=1|x) \geq P(y=0|x)$

$P(y=1|x) \geq P(y=0|x)$

$P(y=1|x) - P(y=1|x) \geq 0$

By Bayes Rule:

$$
\frac{P(x|y=1)P(y=1)}{\sum_{a\in Val(y)}P(x|y=a)P(y=a)} - \frac{P(x|y=0)P(y=0)}{\sum_{a\in Val(y)}P(x|y=a)P(y=a)} \geq 0
$$

Substituting in the equations from the beginning of the problem (and combining the denominator):

$$
\frac{\phi_y\prod_{j=1}^{n}(\phi_{j|y=1})^{x_j}(1-\phi_{j|y=1})^{1-x_j} - (1-\phi_y)\prod_{j=1}^{n}(\phi_{j|y=0})^{x_j}(1-\phi_{j|y=0})^{1-x_j}}{\phi_y\prod_{j=1}^{n}(\phi_{j|y=1})^{x_j}(1-\phi_{j|y=1})^{1-x_j} + (1-\phi_y)\prod_{j=1}^{n}(\phi_{j|y=0})^{x_j}(1-\phi_{j|y=0})^{1-x_j}} \geq 0
$$

$$
\begin{aligned}
&\phi_y\prod_{j=1}^{n}(\phi_{j|y=1})^{x_j}(1-\phi_{j|y=1})^{1-x_j} - (1-\phi_y)\prod_{j=1}^{n}(\phi_{j|y=0})^{x_j}(1-\phi_{j|y=0})^{1-x_j} \geq 0 \\
&\phi_y\prod_{j=1}^{n}(\phi_{j|y=1})^{x_j}(1-\phi_{j|y=1})^{1-x_j} \geq (1-\phi_y)\prod_{j=1}^{n}(\phi_{j|y=0})^{x_j}(1-\phi_{j|y=0})^{1-x_j} \\
&\log\phi_y + \sum_{j=1}^{n}x_j\log\phi_{j|y=1} + \sum_{j=1}^{n}(1-x_j)\log(1-\phi_{j|y=1}) \geq \log(1-\phi_y) + \sum_{j=1}^{n}x_j\log\phi_{j|y=0} + \sum_{j=1}^{n}(1-x_j)\log(1-\phi_{j|y=0})
\end{aligned}
$$

$$\log\phi_y - \log(1-\phi_y) + \sum_{j=1}^{n} x_j \log\phi_{j|y=1} + \sum_{j=1}^{n} \log(1-\phi_{j|y=1}) - \sum_{j=1}^{n} x_j \log(1-\phi_{j|y=1})$$
$$- \sum_{j=1}^{n} x_j \log\phi_{j|y=0} - \sum_{j=1}^{n} \log(1-\phi_{j|y=0}) + \sum_{j=1}^{n} x_j \log(1-\phi_{j|y=0}) \geq 0$$

$$\log\phi_y - \log(1-\phi_y) + \sum_{j=1}^{n} \log(1-\phi_{j|y=1}) - \sum_{j=1}^{n} \log(1-\phi_{j|y=0})$$
$$+ \sum_{j=1}^{n} x_j[\log\phi_{j|y=1} - \log(1-\phi_{j|y=1}) - \log\phi_{j|y=0} + \log(1-\phi_{j|y=0})] \geq 0$$

define $\theta \in \mathbb{R}^{n+1}$ as $\begin{bmatrix}\theta_0\\ \theta_1\\ \vdots\\ \theta_n\end{bmatrix}$

let $\theta_0 = \log\phi_y - \log(1-\phi_y) + \sum_{j=1}^{n} \log(1-\phi_{j|y=1}) - \sum_{j=1}^{n} \log(1-\phi_{j|y=0})$
let $\theta_1 = \log\phi_{1|y=1} - \log(1-\phi_{1|y=1}) - \log\phi_{1|y=0} + \log(1-\phi_{1|y=0})$
let $\theta_2 = \log\phi_{2|y=1} - \log(1-\phi_{2|y=1}) - \log\phi_{2|y=0} + \log(1-\phi_{2|y=0})$
...
let $\theta_n = \log\phi_{n|y=1} - \log(1-\phi_{n|y=1}) - \log\phi_{n|y=0} + \log(1-\phi_{n|y=0})$

then the above inequality can be rewritten as

$\theta_0 + \theta_1 x_1 + \theta_2 x_2 + ... + \theta_n x_n \geq 0$

$\theta^T \begin{bmatrix}1\\ x\end{bmatrix} \geq 0$

5

a)

$$\begin{aligned} P(y;\phi) &= (1-\phi)^{y-1}\phi \\ &= \exp(\log[(1-\phi)^{y-1}\phi]) \\ &= \exp((y-1)\log(1-\phi) + \log\phi) \\ &= \exp(\log(1-\phi)y - \log(1-\phi) + \log\phi) \\ &= \exp(\log(1-\phi)y - (\log(1-\phi) - \log\phi) \\ &= \exp(\log(1-\phi)y - \log(\frac{1-\phi}{\phi})) \end{aligned}$$

We can see this is exponential family with
$\eta = \log(1-\phi)$
$T(y) = y$
$a(\eta) = \log(\frac{1-\phi}{\phi})$

$b(y) = 1$

Solving $\eta = \log(1-\phi)$ for ϕ and substituting into $a(\eta)$, we can find $a(\eta)$ in terms of η:

$$\begin{aligned} \eta &= \log(1-\phi) \\ e^{\eta} &= 1-\phi \\ \phi &= 1-e^{\eta} \end{aligned}$$

$$\begin{aligned} a(\eta) &= \log \frac{1-(1-e^{\eta})}{1-e^{\eta}} \\ a(\eta) &= \log \frac{e^{\eta}}{1-e^{\eta}} \end{aligned}$$

b)

$$\begin{aligned} g(\eta) &= E[T(y);\eta] \\ &= E[y] \\ &= \frac{1}{\phi} \\ &= \frac{1}{1-e^{\eta}} \end{aligned}$$

c)

Since η is linearly related to x, $\eta = \theta^T x$ for some $\theta \in \mathbb{R}^n$ and $\phi = 1 - e^{\theta^T x}$.

$$\begin{aligned} P(y^{(i)}|x^{(i)};\phi) &= \exp(\log(1-\phi)y^{(i)} - \log(\frac{1-\phi}{\phi})) \\ P(y^{(i)}|x^{(i)};\theta) &= \exp(\log(1-(1-e^{\theta^T x^{(i)}}))y^{(i)} - \log(\frac{1-(1-e^{\theta^T x^{(i)}})}{1-e^{\theta^T x^{(i)}}})) \\ P(y^{(i)}|x^{(i)};\theta) &= \exp(\log(e^{\theta^T x^{(i)}})y^{(i)} - \log(\frac{e^{\theta^T x^{(i)}}}{1-e^{\theta^T x^{(i)}}})) \\ \log P(y^{(i)}|x^{(i)};\theta) &= \log(e^{\theta^T x^{(i)}})y^{(i)} - \log(\frac{e^{\theta^T x^{(i)}}}{1-e^{\theta^T x^{(i)}}}) \\ &= \theta^T x^{(i)} y^{(i)} - (\log e^{\theta^T x^{(i)}} - \log(1-e^{\theta^T x^{(i)}})) \\ &= \theta^T x^{(i)} y^{(i)} - \theta^T x^{(i)} + \log(1-e^{\theta^T x^{(i)}}) \end{aligned}$$

$$\begin{aligned}
\frac{\partial}{\partial \theta_j} \log P(y^{(i)}|x^{(i)};\theta) &= x_j^{(i)} y^{(i)} - x_j^{(i)} - \frac{x_j^{(i)} e^{\theta^T x^{(i)}}}{1 - e^{\theta^T x^{(i)}}} \\
&= x_j^{(i)} y^{(i)} - \frac{x_j^{(i)}(1 - e^{\theta^T x^{(i)}}) + x_j^{(i)} e^{\theta^T x^{(i)}}}{1 - e^{\theta^T x^{(i)}}} \\
&= x_j^{(i)} y^{(i)} - \frac{x_j^{(i)}}{1 - e^{\theta^T x^{(i)}}} \\
&= (y^{(i)} - \frac{1}{1 - e^{\theta^T x^{(i)}}}) x_j^{(i)}
\end{aligned}$$

So the stochastic gradient ascent update rule is $\theta_j := \theta_j + \alpha(y^{(i)} - \frac{1}{1 - e^{\theta^T x^{(i)}}}) x_j^{(i)}$